

Computer Sciences Department  
University of Wisconsin-Madison  
CS/ECE 552 – Introduction to Computer Architecture  
In-Class Exercise (03/12) **SOLUTION**

**Answers to all questions should be uploaded on Canvas.**

1. [1 point] You are building a computer system around a processor with in-order execution that runs at 1 GHz and has a CPI of 1, excluding memory accesses. The only instructions that read or write data from/to memory are loads (20% of all instructions) and stores (5% of all instructions).

The memory system for this computer has a split L1 cache. Both the I-cache and the D-cache hold 32 KB each. The I-cache has a 2% miss rate and 64-byte blocks, and the D-cache has a 5% miss rate and 64-byte blocks. The hit time for both the I-cache and the D-cache is 1 ns.

The L2 cache is a unified cache with a total size of 512 KB and a block size of 64-bytes. The hit time of the L2 cache is 15ns for both read hits and write hits. Tag comparison for hit/miss is included in the 15ns in all cases, do not add hit time to miss time on a miss. The local hit rate of the L2 cache is 80%. On a L2 miss, the penalty for both reads and writes is 100 cycles.

Calculate the value of the AMAT *only for data reads* in ns (nanoseconds).

**Solution:**

The formula for AMAT here is:

$$\text{AMAT} = \text{hit time}_{L1} + \text{miss rate}_{L1} \times (\text{hit time}_{L2} + \text{miss rate}_{L2} \times \text{miss penalty}_{L2})$$

From the above, we know:

$$\text{hit time}_{L1} = 1 \text{ processor cycle} = 1 \text{ ns}, \text{ miss rate}_{L1} = 0.05, \text{ hit time}_{L2} = 15 \text{ ns}, \text{ miss rate}_{L2} = 0.80, \text{ miss penalty}_{L2} = 100$$

Thus:

$$\text{AMAT} = 1 + 0.05 \times (15 + 0.8 \times 100) = 2.75 \text{ ns}$$

2. [4 points] Assume you have a direct-mapped cache that uses 8-bit addresses in nibble notation (i.e., base 4), a 32B cache, 4B blocks with the initial cache contents as follows in the *Cache Contents* column and that the initial blocks are accessed in increasing order, and the memory accesses as follows in the *Address* column. Fill in the remaining rows of the table for *Cache Contents*, and in the *Outcome* column specify whether the outcome is a Hit or Miss. When the access is a miss, specify whether the miss is compulsory, capacity, or conflict. We have filled in the first row for you.

**Solution:**

Cache Contents	Address	Outcome
0000, 0010, 0020, 0030, 0100, 0110, 0120, 0130	3020	Miss (compulsory)
0000, 0010, <b>3020</b> , 0030, 0100, 0110, 0120, 0130	3030	Miss (compulsory)
0000, 0010, 3020, <b>3030</b> , 0100, 0110, 0120, 0130	2100	Miss (compulsory)
0000, 0010, 3020, 3030, <b>2100</b> , 0110, 0120, 0130	0012	Hit
0000, 0010, 3020, 3030, 2100, 0110, 0120, 0130	0020	Miss (conflict)
0000, 0010, <b>0020</b> , 3030, 2100, 0110, 0120, 0130	0030	Miss (conflict)
0000, 0010, 0020, <b>0030</b> , 2100, 0110, 0120, 0130	0110	Hit
0000, 0010, 0020, 0030, 2100, 0110, 0120, 0130	0100	Miss (conflict)
0000, 1010, 0020, 0030, <b>0100</b> , 0110, 0120, 0130	2100	Miss (conflict)
1000, 1010, 0020, 0030, <b>2100</b> , 0110, 0120, 0130	3020	Miss (conflict)

3. [5 points] Assume a 32-bit memory address space and a 4-way set-associative cache with an 8-bit set index and 8-bit offset, defined as follows:

```
assign tag[15:0] = address[31:16];
assign set[7:0] = address[15:8];
assign offset[7:0] = address[7:0];
```

Each entry in the table below describes a potential bug in the cache indexing (all lines in the table are Verilog). Specify which of the bugs would still yield a correct cache design; **each bug is applied individually after the assign statements above**. For example, the first bug yields **set[7] = address[15] | address[31]** and **set[6:0] = address[14:8]**. We define a correct cache as satisfying both of the following requirements: 1) a block with address X must not be able to exist at multiple locations in the cache at any one time; and 2) a block with address X must not be confused with a different block with address Y. If a bug yields an incorrect design, give an example of how it breaks one (or both) of the above requirements.

**Solution:**

Bug	Correct?	If incorrect, give example
set[7]  = address[31]	No	Can't distinguish between adds 0x80000000 and 0x80008000
set[7] ^= address[31]	Yes	
tag[0] &= address[15]	No	Can't distinguish between adds 0x80010000 and 0x80000000
tag[0] ^= address[15]	Yes	
set[7] &= address[2]	No	Bytes 0x10008000 and 0x10008004 (same block) map to different sets
set[7] ^= address[2]	No	Bytes 0x10000000 and 0x10000004 (same block) map to different sets
set[0]  = address[9]	No	Can't distinguish between addrs 0x80000200 and 0x80000300
set[0] ^= address[9]	Yes	
tag[15]  = address[16]	No	Can't distinguish between addrs 0x00010000 and 0x80010000
tag[15] ^= address[16]	Yes	
tag[0] &= address[16]	Yes	
tag[0] ^= address[16]	No	Can't distinguish between addrs 0x80010000 and 0x80000000